

Implementation and Evaluation of KNN and SVM Classification Configurations, with Instance Reduction Techniques

Zachary Parent, Sheena Lang, Kacper Poniowski and Carlos Jiménez

November 22, 2024

Contents

1	Introduction	2
1.1	Background	2
1.2	Research Objectives	2
1.3	Methodology Overview	2
2	Data	3
2.1	Dataset Selection and Characteristics	3
2.1.1	Dataset Selection	3
2.1.2	Selection Criteria	3
2.1.3	Dataset Characteristics	4
2.2	Data Preprocessing	4
2.2.1	Handling Different Data Types and Ranges	4
2.2.2	Missing Value Treatment	5
3	Methods	6
4	Results and Analysis	7
5	Conclusion	8
5.1	Key Findings	8
5.2	Practical Implications	8
5.3	Limitations and Future Work	8
5.4	Final Remarks	8
6	Appendix	10

1 Introduction

1.1 Background

1.2 Research Objectives

- 1.

1.3 Methodology Overview

-

2 Data

This section describes the datasets used in our study and details the preprocessing steps applied to prepare them for analysis.

2.1 Dataset Selection and Characteristics

This section outlines our dataset selection criteria and analyzes the characteristics of the chosen datasets.

2.1.1 Dataset Selection

In this project, we aimed to select two datasets that offer substantial variability across different aspects to evaluate the performance of Support Vector Machine (SVM) and k-Nearest Neighbors (KNN) algorithms. The criteria for dataset selection were centered around having one dataset that is small and another that is large, allowing us to analyze both the efficiency and effectiveness of these models across differing dataset sizes.

Smaller datasets

Larger datasets

Attribute types We wanted to assess the models' ability to handle datasets with different types of attributes. For this purpose, we wanted to select one dataset with primarily nominal attributes and another with numerical features. This allows us to evaluate how well each algorithm handles the representation and processing of different data types.

Class distribution By choosing one dataset with a balanced distribution and another with a notable class imbalance, we aim to observe how SVM and KNN, particularly with reduction, perform in scenarios where the data is skewed.

Missing data Lastly, we prioritized finding datasets with varying levels of missing data to examine how effectively the algorithms manage incomplete information.

Based on these criteria, we selected the Hepatitis and Mushroom datasets, as they provide the most distinct and complementary combination for our evaluation.

2.1.2 Selection Criteria

We selected two datasets with contrasting characteristics to evaluate the performance of Support Vector Machine (SVM) and k-Nearest Neighbors (KNN) algorithms across different scenarios. Our selection criteria focused on:

- Dataset size variation (small vs. large)
- Attribute type diversity (nominal vs. numerical)
- Class distribution (balanced vs. imbalanced)
- Missing data patterns

A smaller dataset enables rapid experimentation and initial algorithm validation, while a larger dataset allows evaluation of reduction methods' effectiveness, particularly for KNN's storage and computational requirements.




Figure 1: Class distribution comparison between datasets

2.1.3 Dataset Characteristics

- **Mushroom Dataset**

- 8,124 instances
- 22 nominal attributes
- Binary classification (edible vs. poisonous)
- Nearly balanced classes

- **Hepatitis Dataset**

- 155 instances
- Mixed attribute types (both nominal and numerical)
- Binary classification (survive vs. die)
- Imbalanced classes (79.35% majority class)
- 6.01% missing values

These datasets provide complementary characteristics for evaluating our algorithms’ performance across different scenarios. The Mushroom dataset challenges the algorithms with its size and categorical nature, while the Hepatitis dataset tests their ability to handle mixed data types and class imbalance.

2.2 Data Preprocessing

This section details the preprocessing steps applied to prepare both datasets for analysis.

2.2.1 Handling Different Data Types and Ranges

To manage the varying types and ranges of attributes in our datasets, we implemented specific preprocessing techniques. For the nominal attributes in both the Mushroom and Hepatitis datasets, we used label encoding. This technique converts categorical values into numerical labels, enabling the algorithms to interpret the data correctly. While we considered one-hot encoding to avoid implying any ordinal relationship among categories, we opted for label encoding due to its simplicity and reduced dimensionality, as one-hot encoding would significantly increase dimensions and lead to a sparse space, making accurate predictions more challenging, particularly for KNN with the Mushroom dataset’s numerous nominal features [?].

For the numerical attributes in the Hepatitis dataset, we applied min-max scaling to rescale the data to a fixed range of $[0, 1]$. This normalization is crucial for distance-based algorithms like KNN and SVM, ensuring all features contribute equally to model performance. We evaluated other scaling methods, such as standardization, but chose min-max scaling for its effectiveness in maintaining the original data distribution [?].

We implemented specific preprocessing techniques based on the characteristics of each dataset:

- **Nominal Attributes**

- Applied label encoding to both datasets
- Chose label encoding over one-hot encoding to prevent dimensionality explosion
- Particularly important for the Mushroom dataset’s numerous categorical features

- **Numerical Attributes**

- Applied min-max scaling to the Hepatitis dataset’s numerical features
- Normalized all values to $[0, 1]$ range
- Essential for distance-based algorithms (KNN and SVM)

2.2.2 Missing Value Treatment

Addressing missing values is a critical step in preparing our datasets for analysis, as they can significantly impact model performance. In the case of the nominal attributes in both the Mushroom and Hepatitis datasets, we opted to impute missing values with the majority class. This method is straightforward and effective for maintaining dataset integrity. However, it can also introduce bias, particularly in the Hepatitis dataset, where the majority class represents 79.35% of instances. Relying on this method may lead to a situation where the imputed values disproportionately favor the majority class, thereby affecting the overall distribution and potentially skewing the results [?].

For the numerical attributes in the Hepatitis dataset, we used the mean of the available data to fill in missing values. This approach preserves the overall data distribution and is easy to implement, but it is not without its drawbacks. The mean can be heavily influenced by outliers, which might distort the data and lead to less accurate predictions. This is especially important in medical datasets, where extreme values may carry significant meaning.

We also considered employing K-Nearest Neighbors (KNN) for imputing missing values, as it could provide a more nuanced approach by considering the nearest data points for each instance. However, we ultimately decided against this option to avoid introducing bias into our evaluation. Since KNN is one of the algorithms we are testing, using it for imputation could influence its performance and lead to skewed results. Therefore, we chose the more straightforward methods of majority class imputation for nominal values and mean imputation for numerical values, allowing for a clearer assessment of the models’ effectiveness without confounding factors.

We employed different strategies for handling missing values based on attribute type:

- **Nominal Attributes**

- Imputed with mode (majority class)
- Applied to both datasets
- Potential limitation: May reinforce majority class bias

- **Numerical Attributes**

- Imputed with mean values
- Applied only to Hepatitis dataset
- Preserves overall distribution while handling missing data

Alternative approaches such as KNN-based imputation were considered but rejected to avoid introducing bias into our evaluation of KNN as a classifier. Our chosen methods provide a balance between simplicity and effectiveness while maintaining data integrity.

3 Methods

This section details our experimental methodology, including algorithm implementations, parameter configurations, and statistical analysis approaches. We present our methods in four main parts:

4 Results and Analysis

5 Conclusion

5.1 Key Findings

-

5.2 Practical Implications

-

5.3 Limitations and Future Work

-

5.4 Final Remarks

References

6 Appendix